

The 3 Core Prompting Frameworks

These aren't tips — they're thinking systems. Master the structure, and every AI tool becomes predictable.

RTF Framework (Role / Task / Format)

R = Role — Assign a specific expert identity with years of experience and domain expertise.

T = Task — Define the exact deliverable with scope, constraints, and success criteria.

F = Format — Specify output structure, length, tone, and any formatting requirements.

CO-STAR Framework (Context / Objective / Style / Tone / Audience / Response)

C = Context — Background situation, constraints, and information the AI needs to reason correctly.

O = Objective — The specific outcome — not just "write" but what the output must achieve.

S = Style — Writing style: academic, conversational, technical, journalistic, Socratic, etc.

T = Tone — Emotional register: authoritative, empathetic, provocative, measured, urgent.

A = Audience — Who reads this? Their knowledge level, concerns, and what motivates them.

R = Response — Exact format: word count, structure, sections, deliverable type.

RASCEF Framework (Role / Action / Steps / Context / Examples / Format)

R = Role — Expert identity with specific credentials that shape reasoning approach.

A = Action — The core task decomposed into its primary objective.

S = Steps — Numbered sub-tasks that force sequential, chain-of-thought reasoning.

C = Context — Hard constraints, edge cases, and rules the AI must respect.

E = Examples — One or more examples of desired output quality and structure.

F = Format — Exact output spec — length, structure, file format, labeling conventions.

RTF Quick tasks with clear scope. Daily driver. 80% of use cases.	CO-STAR Content + communication where audience and tone matter more than steps.	RASCEF Complex, multi-step work. When you need the AI to reason, not just generate.
--	--	--

Advanced Techniques Used in This Guide: [\[Chain-of-Thought\]](#) = force step-by-step reasoning | [\[Self-Evaluation\]](#) = AI critiques its own output | [\[Constraint Stacking\]](#) = layered rules that eliminate generic output | [\[Persona Depth\]](#) = detailed role with reasoning style | [\[Output Calibration\]](#) = explicit quality benchmarks | [\[Few-Shot\]](#) = examples that set the pattern

15 Advanced Prompt Templates

Each prompt uses multiple techniques stacked together. Swap [variables] with your specifics. Notice the structure — that's the logic.

MARKETING & GROWTH

1. High-Converting Email Sequence Architect [Chain-of-Thought] [Constraint Stacking] [Self-Evaluation]

You are a direct-response email strategist who has generated \$10M+ in email revenue for SaaS companies. Your task: design a 5-email welcome sequence for [product/service] targeting [specific audience with their core frustration]. Before writing, reason through these steps internally: (1) Map the psychological journey from skeptic to buyer — what objection does each email need to dissolve? (2) Identify the single emotional trigger per email that drives the open and the click. (3) Plan the value escalation — each email must deliver more value than the last while increasing commitment. Now write all 5 emails. For each email include: subject line (under 45 chars, no clickbait), preview text (under 90 chars), full body copy, CTA button text, and send timing relative to signup. Constraints: Email 1 must deliver immediate value with zero selling. Email 3 is the pivot from pure value to soft pitch. Email 5 is the direct ask with a deadline. After writing all 5, critique your own sequence: identify the weakest email and rewrite it stronger.

2. Audience-Aware Ad Copy Matrix [Persona Depth] [Constraint Stacking] [Output Calibration]

You are a performance marketer who manages \$500K/month in ad spend across Meta, Google, and LinkedIn. You understand that the same product needs radically different messaging for different awareness levels. Product: [product/service]. Create an ad copy matrix with 3 variations for EACH of these 3 audience segments: (A) Problem-Unaware — they don't know they have the problem yet. (B) Problem-Aware — they know the pain but not the solution. (C) Solution-Aware — they know solutions exist but haven't chosen yours. Each variation must include: primary text (max 125 chars), headline (max 40 chars), description (max 30 chars), and the psychological principle it leverages (curiosity gap, social proof, loss aversion, etc.). Format as a table with columns: Segment | Variation | Primary Text | Headline | Description | Principle. Quality benchmark: every headline must pass this test — would a person scrolling at 3am stop and read it?

3. Viral Content Reverse-Engineer [Chain-of-Thought] [Few-Shot] [Self-Evaluation]

You are a content strategist who studies viral mechanics across platforms. I will give you a piece of content that performed well: [paste the content or describe it]. Your task: (1) Deconstruct WHY it worked — analyze the hook structure, emotional trigger, information gap, shareability mechanics, and format advantages. (2) Extract the reusable framework — strip away the specific topic and identify the structural template underneath. (3) Generate 5 new content ideas in my niche [your niche] using that same structural template, but with original angles. For each idea, include: the hook (first line), the content structure, the emotional payoff, and the CTA. (4) Rank your 5 ideas from highest to lowest viral potential and justify each ranking with one sentence explaining the trade-off between reach and depth.

WRITING & CONTENT STRATEGY

4. Authority-Building Article Architect [Chain-of-Thought] [Persona Depth] [Output Calibration]

You are an editorial strategist for a top-tier publication like Harvard Business Review. You specialize in turning domain expertise into authoritative thought leadership. My expertise: [your area of expertise]. My target reader: [describe your ideal reader — their role, knowledge level, and what keeps them up at night]. Step 1: Generate 5 article concepts that position me as the definitive expert in this space. Each concept must have a contrarian or non-obvious angle — avoid anything a reader would respond to with 'obviously.' Step 2: For the strongest concept, build a full outline: compelling title (use a specific number, question, or provocation), opening hook that creates an information gap in the first 2 sentences, 4-5 H2 sections each building on the previous one's conclusion, one data point or case study per section (suggest what type of evidence to include), and a conclusion that reframes the reader's understanding of the topic. Quality benchmark: every H2 heading should work as a standalone social media post.

5. Voice Clone & Style Transfer Engine [Few-Shot] [Constraint Stacking] [Output Calibration]

Your task is to analyze my writing style and then replicate it precisely. Here are 3 samples of my writing: [paste sample 1] [paste sample 2] [paste sample 3]. Step 1 – Style Analysis: Before writing anything, produce a style profile covering: average sentence length, vocabulary complexity (simple/moderate/advanced), use of metaphors and analogies (frequency and type), paragraph rhythm (short-long patterns), tone markers (formal/casual/hybrid), signature phrases or patterns, and how I handle transitions between ideas. Step 2 – Calibration Test: Write a single paragraph about [any topic I haven't written about] in my exact voice. Step 3 – Flag any elements of my style that could be improved for clarity or engagement without losing my voice, with specific examples of how each change would read. Constraint: do NOT homogenize my style into 'standard AI output.' Preserve the quirks – they are the voice.

15 Advanced Prompt Templates (continued)

6. Adaptive Newsletter System [Persona Depth] [Chain-of-Thought] [Constraint Stacking]

You are a newsletter strategist who has grown 3 newsletters past 50K subscribers. My niche: [your niche]. My newsletter frequency: [weekly/biweekly]. My brand voice: [describe in 2-3 adjectives]. Design a repeatable newsletter template system with these constraints: (1) The structure must work every single week without feeling repetitive to readers. (2) Each issue must contain exactly 3 sections that serve different purposes – one for new subscriber retention, one for engaged reader deepening, one for revenue (affiliate, product, sponsor). (3) The opening must use a different hook type each week – rotate between: personal story, contrarian take, data point, reader question, prediction. Deliver: the full template with placeholder copy for each section, 5 weeks of specific issue plans with topics and hook types, and a subject line formula that averages 40%+ open rates in [your niche] with 3 example subject lines per issue.

CODING & TECHNICAL PROBLEM-SOLVING

7. Senior Code Reviewer with Architectural Thinking [Persona Depth] [Chain-of-Thought] [Self-Evaluation]

You are a staff-level [language] engineer who reviews pull requests at a company that processes 10M+ requests/day. You think about code in three layers: correctness (does it work?), maintainability (will a junior dev understand this in 6 months?), and scalability (what breaks when traffic 10x's?). Review this code: [paste code]. Produce your review in this exact structure: (1) CRITICAL – bugs, security issues, or data loss risks. Must fix before merge. (2) IMPORTANT – performance bottlenecks, missing error handling, or architectural concerns. Should fix. (3) SUGGESTIONS – readability improvements, naming conventions, design pattern opportunities. Nice to have. (4) WHAT'S GOOD – explicitly call out 2-3 things done well (this matters for team morale). (5) REFACTORED VERSION – provide the complete corrected code with inline comments explaining every change. Constraint: assume this code runs in production. Flag any operation that could fail silently.

8. Architecture Decision Record Generator [Chain-of-Thought] [Constraint Stacking] [Output Calibration]

You are a solutions architect evaluating technology choices. I need to decide: [describe the technical decision, e.g., 'database choice for a real-time analytics dashboard' or 'monolith vs microservices for our MVP']. Context: [team size], [current stack], [expected scale in 12 months], [budget constraints]. Produce an Architecture Decision Record (ADR) with this structure: (1) Decision Title (2) Status: Proposed (3) Context – what forces are at play (business, technical, team) (4) Options Considered – minimum 3 options, each with: description, pros (at least 3), cons (at least 3), cost estimate (time + money), and hidden risk most teams overlook (5) Decision – your recommended option with reasoning chain showing how you weighed the trade-offs (6) Consequences – what this decision makes easier AND what it makes harder in 12 months (7) Reversal Cost – how expensive is it to change this decision later (low/medium/high and why). Constraint: do NOT default to the most popular option. Reason from my specific context.

9. Debug Detective with Root Cause Analysis [Chain-of-Thought] [Self-Evaluation] [Constraint Stacking]

You are a senior debugger who has resolved thousands of production incidents. Your debugging methodology: reproduce, isolate, hypothesize, verify, fix, prevent. Here is my problem: [describe the bug – what you expected, what actually happened, error messages if any, and what you have already tried]. Code: [paste relevant code]. Do NOT jump to a solution. Instead: (1) Form 3 hypotheses ranked by probability, explaining your reasoning for each. (2) For the most likely hypothesis, describe exactly how you would confirm it – what would you log, print, or test? (3) Provide the fix with a clear explanation of why the bug occurred at a conceptual level (not just 'change line 42'). (4) Write a regression test that would have caught this bug before it shipped. (5) Identify one defensive coding pattern that would prevent this entire category of bug in the future. If my description is missing information you need, tell me exactly what to check instead of guessing.

BUSINESS STRATEGY & ANALYSIS

10. Pre-Mortem Risk Analyst [Chain-of-Thought] [Persona Depth] [Constraint Stacking]

You are a strategic risk consultant who specializes in pre-mortem analysis – imagining a project has already failed and working backward to understand why. My project: [describe the project, timeline, team, and goals]. Conduct a pre-mortem: (1) Imagine it is [deadline + 2 months] and this project has completely failed. Generate the 7 most likely causes of failure, ranked by probability. For each cause: describe the failure scenario in one paragraph, identify the earliest warning sign that would have been visible, and propose one specific preventive action that costs less than 4 hours to implement. (2) Now identify 2 'black swan' failure modes – unlikely but catastrophic risks that most teams never consider. (3) Create a one-page risk dashboard I can review weekly: a table with columns for Risk, Likelihood (1-5), Impact (1-5), Early Warning Signal, Owner, and Mitigation Status. Constraint: at least 3 of your failure causes must be people/communication problems, not just technical risks.

11. Decision Matrix Builder with Hidden Assumptions [Chain-of-Thought] [Self-Evaluation] [Output Calibration]

You are a management consultant trained in structured decision-making. I need to decide between: [Option A] vs [Option B] vs [Option C]. Context: [describe your situation, constraints, and what matters most to you]. Build a weighted decision matrix, but first: (1) Before scoring anything, surface 5 assumptions I am probably making unconsciously about this decision. For each assumption, explain how the decision changes if that assumption is wrong. (2) Propose 8 evaluation criteria – 5 that I would naturally think of, and 3 that most people overlook in this type of decision. (3) Assign weights to each criterion based on my stated context (weights must sum to 100%). Show your reasoning for the weighting. (4) Score each option 1-10 on each criterion with a one-line justification per score. (5) Calculate weighted totals and declare a winner. (6) Then argue AGAINST your own recommendation – what is the strongest case for not choosing the winner? Format: structured sections, final recommendation in bold with confidence level (high/medium/low).

15 Advanced Prompt Templates (continued)

12. Competitor Teardown Analyst [Chain-of-Thought] [Constraint Stacking] [Output Calibration]

You are a competitive intelligence analyst. Target competitor: [name]. My company/product: [name + brief description]. Produce a teardown structured as: (1) Positioning Map – how do they position themselves vs the market? What words do they use repeatedly? What do they deliberately avoid saying? (2) Strengths to Steal – 3 things they do better than us that we could realistically adopt within 90 days. (3) Weaknesses to Exploit – 3 gaps in their offering, messaging, or audience coverage that represent our opportunity. (4) Their Next Move – based on their recent trajectory, predict their most likely strategic move in the next 6 months. (5) One-Line Battle Card – a single sentence our sales/marketing team can use when a prospect says 'but what about [competitor]?' Constraint: distinguish between publicly verifiable facts and your inferences. Label each clearly.

THINKING & PRODUCTIVITY SYSTEMS

13. Feynman Learning Accelerator [Chain-of-Thought] [Self-Evaluation] [Output Calibration]

You are a learning scientist who applies the Feynman Technique: if you cannot explain it simply, you do not understand it. I want to deeply understand [topic]. My current level: [beginner/intermediate/advanced]. Execute this sequence: (1) Explain [topic] as if I am a smart 12-year-old. Use one analogy from everyday life. No jargon. (2) Now explain it at a college level, introducing the proper terminology – but connect each technical term back to the simple analogy from step 1 so I see how they map. (3) Identify the 3 most common misconceptions people have about this topic and explain why each is wrong. (4) Give me 3 questions I should be able to answer if I truly understand this. Do NOT give me the answers yet – let me try first. (5) After I respond with my answers, evaluate them and fill in any gaps. Constraint: if at any point your explanation requires a concept I might not know, pause and explain that concept first.

14. Second-Order Thinking Simulator [Chain-of-Thought] [Persona Depth] [Constraint Stacking]

You are a systems thinker trained in second and third-order consequence analysis. I am considering this action/decision: [describe the decision]. Map the consequences across three orders: FIRST-ORDER (immediate, obvious): What happens right away? List 3-5 direct consequences. SECOND-ORDER (6-12 months, non-obvious): What do those first-order effects cause? Who adapts, and how? What incentives shift? List 3-5 second-order effects that most people fail to anticipate. THIRD-ORDER (12-36 months, systemic): What new equilibrium emerges? What unintended consequences crystallize into permanent structural changes? List 2-3 third-order effects. Then: (1) Identify which order of consequence poses the highest risk to my goals. (2) Suggest one modification to my original decision that preserves the first-order benefits while mitigating the worst second-order risk. Format: visual cascade – indent each order further right to show the chain.

15. Weekly Strategic Review System [Persona Depth] [Constraint Stacking] [Output Calibration]

You are my strategic thinking partner – a chief of staff who asks hard questions. Context about my current situation: [paste a brief summary: what you are working on, your top 3 goals, what happened this week]. Run this weekly review protocol: (1) WINS – Identify the most strategically important thing I accomplished this week (not just the busiest). (2) MISALIGNMENT CHECK – Based on my stated goals, flag anything I spent significant time on that does NOT directly advance a top-3 goal. Ask me to justify it or cut it. (3) BOTTLENECK – What is the single biggest constraint slowing my progress right now? Categorize it: is it a knowledge gap, a resource constraint, a decision I am avoiding, or a dependency on someone else? (4) NEXT WEEK'S ONE THING – If I could only accomplish one thing next week that would create the most forward momentum, what should it be and why? (5) HARD QUESTION – Ask me one uncomfortable question about my strategy that I am probably avoiding. Constraint: be direct. Do not soften feedback with qualifiers. I need a thought partner, not a cheerleader.

Before vs. After: The Logic Difference

Same intent, radically different structure. The techniques in [brackets] show you what's doing the heavy lifting.

Scenario: Writing a marketing email for a product launch [Techniques: Persona Depth, Constraint Stacking, Self-Evaluation]

BEFORE — The way most people prompt

Write me a marketing email for my new product.

AFTER — The Logic of Logic way

You are a direct-response copywriter who has generated \$10M+ in email revenue for SaaS companies. Write a launch email for my new AI prompt course targeting small business owners who have tried ChatGPT but felt overwhelmed by the blank screen. Structure: open with a 1-sentence relatable pain point they will nod along to, introduce the course as the answer to that specific pain, include 3 bullet-point benefits (each starting with a verb), and end with a single CTA button to [URL] with urgency tied to a launch-week bonus. Tone: confident but warm – like a smart friend who just figured something out. Length: 250 words max. Also provide 3 subject line options under 45 characters each. After writing, identify the weakest sentence in the email and rewrite it stronger.

Why it works: The vague prompt gives zero context — the AI guesses your product, audience, voice, and format. The logical prompt uses a credentialed persona, specific audience psychology, structural constraints, and a self-evaluation loop. The AI cannot produce generic output because every constraint eliminates a lazy shortcut.

Scenario: Debugging code that isn't working [Techniques: Chain-of-Thought, Constraint Stacking, Self-Evaluation]

BEFORE — The way most people prompt

Fix my Python code it's not working.

AFTER — The Logic of Logic way

You are a senior Python engineer who debugs by hypothesis, not by guessing. My script reads a CSV and calculates monthly revenue totals. Expected: grouped totals by month. Actual: KeyError on 'date' column, even though the column exists in the file. I have already tried: stripping whitespace from headers, and confirming the file loads with pandas.read_csv(). Code: [paste code]. CSV columns: date (YYYY-MM-DD), product (string), amount (float). Do NOT jump to a fix. Instead: (1) State 3 hypotheses ranked by likelihood. (2) Provide the fix with a conceptual explanation of why the bug occurs. (3) Write one test case that would catch this bug in CI. (4) Name the defensive pattern that prevents this bug category entirely.

Why it works: The vague prompt forces the AI to hallucinate context. The logical prompt provides the error, what was already tried, the data schema, and forces hypothesis-first debugging instead of 'try this' shotgun fixes. Result: an accurate root-cause analysis on the first try instead of 5 rounds of back-and-forth.

Scenario: Learning a complex new topic [Techniques: Feynman Technique, Chain-of-Thought, Constraint Stacking]

BEFORE — The way most people prompt

Teach me about machine learning.

AFTER — The Logic of Logic way

You are a learning scientist who applies the Feynman Technique. I am a marketing manager with zero coding experience. My goal: understand ML well enough to have informed conversations with my company's data team within 2 weeks. I do NOT need to build models – I need to understand what is possible, what is hype, and what questions to ask. Step 1: Explain ML as if I am a smart 12-year-old, using a cooking analogy. Step 2: Re-explain with proper terminology, mapping each technical term back to the cooking analogy. Step 3: List the 3 most common misconceptions non-technical leaders have about ML and why each is wrong. Step 4: Give me a 2-week learning plan: one 20-minute task per day, free resources only, with a self-check quiz at the end of each week. Step 5: Write me 5 questions I can ask my data team next Monday that will make them take me seriously. Constraint: if any explanation requires a concept I might not know, explain that concept first before proceeding.

Why it works: The vague prompt produces a textbook overview. The logical prompt defines your exact starting point, specific goal, timeline, non-goal (not building models), preferred analogy style, and output format. The AI now builds a personalized curriculum anchored to your real situation instead of a Wikipedia summary with no actionable next steps.

This cheat sheet is your starting framework. The real skill is learning to stack these techniques instinctively — seeing a blank prompt box and automatically thinking: what role, what constraints, what format, what self-check? That is the Logic of Logic. Visit logicoflogic.com for weekly breakdowns of advanced prompting techniques, real-world case studies, and a community of people who think before they prompt.